

Screenshot of the connection (i.e. showing your terminal/command-line information):

```
Cloud Shell
Terminal (cs411-team004-orange) x + -
Open Editor

Welcome to Cloud Shell! Type "help" to get started.
Your Cloud Platform project in this session is set to cs411-team004-orange.
Use "gcloud config set project [PROJECT_ID]" to change to a different project.
rohanvij2@cloudshell:~ (cs411-team004-orange)$ gcloud sql connect cs411-team004-orange-db --user=root
Allowlisting your IP for incoming connection for 5 minutes...done.
Connecting to database with SQL user [root]. Enter password:
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 25
Server version: 8.0.26-google (Google)

Copyright (c) 2000, 2023, Oracle and/or its affiliates.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql>
```

Provided the DDL commands for your tables (-0.5% for each mistake):

```
CREATE TABLE Athletes
(
    Ath_Name VARCHAR(512) PRIMARY KEY,
    NOC VARCHAR(512),
    Discipline VARCHAR(512) REFERENCES Sport(Discipline),
    Age INT,
    Salary INT
);
```

```
CREATE TABLE Coaches
(
    Coa_Name VARCHAR(512) PRIMARY KEY,
    NOC VARCHAR(512),
    Discipline VARCHAR(512),
    Event VARCHAR(512)
);
```

```
CREATE TABLE Sport
(
    Discipline VARCHAR(512) PRIMARY KEY,
    Avg_Salary INT,
    Avg_Age INT
);
```

```
CREATE TABLE Teams
(
    Name CHAR(100),
    Discipline CHAR(100),
    NOC CHAR(100),
    Event VARCHAR(512),
```

```
PRIMARY KEY (NOC, Discipline, Event),  
FOREIGN KEY (Discipline) REFERENCES Sport(Discipline) ON DELETE CASCADE  
);
```

```
CREATE TABLE Medals (  
  Ath_Name CHAR(100),  
  NOC CHAR(100),  
  Num_total INT,  
  Num_Gold INT,  
  Num_Silver INT,  
  Num_Bronze INT,  
  PRIMARY KEY (Ath_Name, NOC),  
  FOREIGN KEY (Ath_Name) REFERENCES Athletes(Ath_Name)  
);
```

```
CREATE TABLE Athletes_Coaches (  
  Ath_Name CHAR(100),  
  Coa_Name CHAR(100),  
  
  PRIMARY KEY (Ath_Name, Coa_Name),  
  FOREIGN KEY (Ath_Name) REFERENCES Athletes(Ath_Name),  
  FOREIGN KEY (Coa_Name) REFERENCES Coaches(Coa_Name)  
);
```

```
CREATE TABLE Team_Coaches (  
  Coa_Name CHAR(100),  
  NOC CHAR(100),  
  Discipline CHAR(100),  
  Event CHAR(100),  
  
  PRIMARY KEY (Coa_Name, NOC, Discipline, Event),  
  
  FOREIGN KEY (Coa_Name) REFERENCES Coaches(Coa_Name),  
  FOREIGN KEY (NOC) REFERENCES Teams(NOC),  
  FOREIGN KEY (Discipline) REFERENCES Teams(Discipline)  
);
```

Inserted at least 1000 rows in the tables

You should upload a screenshot of a count query to show this, 1% for each table

You should insert at least 1000 rows each in three of the tables:

```
CLOUD SHELL
Terminal (cs411-team004-orange) × + ▾

| information_schema |
| mysql              |
| olympicstokyo      |
| performance_schema |
| sys                |
+-----+
5 rows in set (0.00 sec)

mysql> use olympicstokyo;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
mysql> show tables;
+-----+
| Tables_in_olympicstokyo |
+-----+
| Athletes                |
| Athletes_Coaches        |
| Coaches                  |
| Medals                   |
| Sport                    |
| Team_Coaches             |
| Teams                    |
+-----+
7 rows in set (0.00 sec)

mysql> SELECT COUNT(*) FROM Athletes;
+-----+
| COUNT(*) |
+-----+
|      11062 |
+-----+
1 row in set (0.01 sec)

mysql> SELECT COUNT(*) FROM Coaches;
+-----+
| COUNT(*) |
+-----+
|       382 |
+-----+
1 row in set (0.00 sec)

mysql> SELECT COUNT(*) FROM Medals;
+-----+
| COUNT(*) |
+-----+
|      5066 |
+-----+
1 row in set (0.00 sec)

mysql> SELECT COUNT(*) FROM Sport;
+-----+
| COUNT(*) |
+-----+
|        46 |
+-----+
1 row in set (0.00 sec)

mysql> SELECT COUNT(*) FROM Teams;
+-----+
| COUNT(*) |
+-----+
|      1030 |
+-----+
1 row in set (0.00 sec)

mysql> 
```

Two advanced queries:

-- Query 1:

```
SELECT DISTINCT Athletes.Ath_Name, Athletes.Discipline, Athletes.NOC, Coaches.Event,
Teams.Name
FROM Athletes
JOIN Coaches ON Athletes.NOC = Coaches.NOC
JOIN Teams ON Teams.NOC = Athletes.NOC AND Teams.Discipline = Athletes.Discipline AND
Teams.Event = Coaches.Event
```

```
WHERE Athletes.Age > 30
AND Athletes.Salary > (SELECT AVG(Avg_Salary) FROM Sport WHERE Discipline =
Athletes.Discipline);
```

-- Query 1 calculates the names, disciplines, NOCs, events, and teams of all athletes over the age of 30 who are making more than the average salary for their discipline.

```
mysql> SELECT DISTINCT Athletes.Ath_Name, Athletes.Discipline, Athletes.NOC, Coaches.Event, Teams.Name
-> FROM Athletes
-> JOIN Coaches ON Athletes.NOC = Coaches.NOC
-> JOIN Teams ON Teams.NOC = Athletes.NOC AND Teams.Discipline = Athletes.Discipline AND Teams.Event = Coaches.Event
-> WHERE Athletes.Age > 30
-> AND Athletes.Salary > (SELECT AVG(Avg_Salary) FROM Sport WHERE Discipline = Athletes.Discipline)
-> LIMIT 15;
```

Ath_Name	Discipline	NOC	Event	Name
ABALDE Alberto	Basketball	Spain	Women	Spain
ABALDE Alberto	Basketball	Spain	Men	Spain
ABBINGH Lois	Handball	Netherlands	Women	Netherlands
ABRAHAMS Rusten John	Hockey	South Africa	Women	South Africa
ABRAHAMS Rusten John	Hockey	South Africa	Men	South Africa
ABRINES REDONDO Alejandro	Basketball	Spain	Women	Spain
ABRINES REDONDO Alejandro	Basketball	Spain	Men	Spain
ACHIUWA Precious	Basketball	Nigeria	Women	Nigeria
ACHIUWA Precious	Basketball	Nigeria	Men	Nigeria
ADEL Ibrahim	Football	Egypt	Men	Egypt
AEDO Yanara	Football	Chile	Women	Chile
AGATSUMA Haruka	Baseball/Softball	Japan	Softball	Japan
AGATSUMA Haruka	Baseball/Softball	Japan	Baseball	Japan
AGUIRRE Eduardo	Football	Mexico	Men	Mexico
AKOZ Simge Sebnem	Volleyball	Turkey	Women	Turkey

15 rows in set (0.01 sec)

Image of the first 15 rows for query 1

```
-- Query 2:
SELECT Coaches.Coa_Name, Coaches.Discipline, Coaches.Event, SUM(Medals.Num_Gold)
AS Total_Gold
FROM Coaches
JOIN Medals ON Coaches.NOC = Medals.NOC
JOIN Teams ON Coaches.NOC = Teams.NOC AND Coaches.Discipline = Teams.Discipline
AND Coaches.Event = Teams.Event
GROUP BY Coaches.Coa_Name, Coaches.Discipline, Coaches.Event
HAVING SUM(Medals.Num_Gold) > 5;
```

-- Query 2: This query selects the coaches' names, disciplines, events, and the total number of gold medals won by teams coached by each coach. It only returns results for coaches who have won at least 5 gold medals.

Coa_Name	Discipline	Event	Total_Gold
TERZIC Zoran	Volleyball	Women	30
SAVIC Dejan	Water Polo	Men	30
MALJKOVIC Marina	Basketball	Women	30
ZAITSEVA Olesia	Artistic Swimming	Duet	72
MEZHENINA Valeriia	Artistic Swimming	Team	72
SACCHETTI Romeo	Basketball	Men	165
PIZZOLINI Federico	Baseball/Softball	Softball	165
PALERMI Giovanna	Baseball/Softball	Softball	165
MAZZANTI Davide	Volleyball	Women	165
GIALLOMBARDO Patrizia	Artistic Swimming	Duet	165
FARINELLI Roberta	Artistic Swimming	Duet	165
CAMPAGNA Alessandro	Water Polo	Men	165
BLENGINI Gianlorenzo	Volleyball	Men	165
THOMADIS Lissa	Basketball	Women	166
SZAUDEK Gabor	Artistic Swimming	Duet	166

15 rows in set (0.10 sec)

EXPLAIN ANALYZE BEFORE INDEXING QUERY 1:

[illegible][illegible]

Before indexing, the running time of query 1 is 0.29 seconds and the running time for query 2 is 0.11 seconds from the images above. So the average time cost is $(0.29+0.11)/2 = 0.2$ second.

We try to add indices to the table that is used by both queries, which is Teams.

[illegible]

At this time, the second query takes 0.1 second which is faster than 0.11 second before, but the first query takes a longer time as well which is 0.31 second. So the average time cost is $(0.31+0.1)/2 = 0.205$ that the performance hasn't changed a lot. So it's still not a good choice.

Index design3

```
CREATE INDEX salary_idx ON Athletes (Salary);
```

```
CREATE INDEX coa_idx ON Coaches(Discipline)/CREATE INDEX coa_idx ON Coaches(NOC);
```

We take a closer look at the two query and identify table Coaches need more indices due to the GROUP BY implementation in query2. We also add index to the Salary in athletes table due to the comparison implementation in query1.

```
--> Table scan on <temporary> (cost=0.01..285.38 rows=22630) (actual time=0.003..0.211 rows=1140 loops=1)
--> Temporary table with deduplication (cost=19230.43..19515.79 rows=22630) (actual time=200.038..200.311 rows=1140 loops=1)
--> Nested loop inner join (cost=16967.40 rows=22630) (actual time=0.219..137.318 rows=4197 loops=1)
--> Nested loop inner join (cost=3046.82 rows=22630) (actual time=0.192..80.526 rows=27803 loops=1)
--> Filter: ((Athletes.Age > 30) and (Athletes.Salary > (select #2)) and (Athletes.NOC is not null) and (Athletes.Discipline is not null)) (cost=1126.25 rows=3673) (actual time=0.084..22.641 rows=2840 loops=1)
--> Table scan on Athletes (cost=1126.25 rows=11020) (actual time=0.039..4.887 rows=11062 loops=1)
--> Select #2 (subquery in conditions depends)
--> Aggregate: avg(Sport.Avg_Salary) (cost=0.45 rows=1) (actual time=0.002..0.002 rows=1 loops=5147)
--> Single-row index lookup on Sport using PRIMARY (Discipline=Athletes.Discipline) (cost=0.35 rows=1) (actual time=0.002..0.002 rows=1 loops=5147)
--> Filter: (Coaches.'Event' is not null) (cost=1.54 rows=6) (actual time=0.010..0.020 rows=10 loops=2840)
--> Index lookup on Coaches using Coaches.noc_idx (NOC=Athletes.NOC) (cost=1.54 rows=6) (actual time=0.010..0.019 rows=10 loops=2840)
--> Filter: ((Teams.NOC = Athletes.NOC) and (Teams.Discipline = Athletes.Discipline) and (Teams.'Event' = Coaches.'Event')) (cost=0.25 rows=1) (actual time=0.004..0.004 rows=0 loops=27803)
--> Single-row index lookup on Teams using PRIMARY (NOC=Athletes.NOC, Discipline=Athletes.Discipline, Event=Coaches.'Event') (cost=0.25 rows=1) (actual time=0.003..0.003 rows=0 loops=27803)
```

[illegible]

The time cost for the first query decreases a lot from 0.29 to 0.2 seconds. But the second query doesn't change the time cost which means adding an index to table Coaches doesn't improve the performance of query2.

In addition, We also try to add an index to the key in the Table Medals. But the time cost of the query is still about 0.1 seconds.

To conclude, we find adding indices to the Athletes table will make query1 perform better. But we can't find a better index to make query2 work better.

The reason why some of the indexing will make databases slower is

1. Indexes can cause overhead: The more indexes you have on a table, the more overhead there is in terms of storage space and processing power required to maintain those indexes. This overhead can slow down queries.

2. Index selection: The query optimizer must select the best index to use for a given query. When there are many indexes on a table, the optimizer has more options to consider, which can increase the time required to choose the best index.